

# MPI Basics with Python

Iman Rahbari, Computation Scientist

Cesar Sul, Research Computing Facilitator

# Outline


- Parallel Computing
- What is MPI
- Running programs in parallel
- Dividing work among workers
- Point to point communication
- Collective communication

# Parallel Computing:

Parallel computing is when multiple *processors* or *computers* work together to solve a problem at the same time.


Main difference between these two cases is the **underlying memory**

Multiple *processors*:  
**Shared Memory**



All access the same memory  
Multi-Processing  
OpenMP programming

Multiple *computers*:  
**Distributed Memory**




Each CPU has a separate memory  
All CPUs communicate via network  
MPI programming

# Parallel Computing:

Parallel computing is when multiple *processors* or *computers* work together to solve a problem at the same time.


Main difference between these two cases is the **underlying memory**

Multiple *processors*:  
Shared Memory




All access the same memory  
Multi-Processing  
OpenMP programming

**Hybrid**



A combination of both  
MPI+OpenMP hybrid programming

Multiple *computers*:  
Distributed Memory




Each CPU has a separate memory  
All CPUs communicate via network  
MPI programming

# Parallel Computing:

Parallel computing is when multiple *processors* or *computers* work together to solve a problem at the same time.


Main difference between these two cases is the **underlying memory**

Multiple *processors*:  
Shared Memory




All access the same memory  
Multi-Processing  
OpenMP programming

**Hybrid**



A combination of both  
MPI+OpenMP hybrid programming

Multiple *computers*:  
Distributed Memory



Each CPU has a separate memory  
All CPUs communicate via network

**MPI programming**  
**This is our focus here!**

# What is MPI?

- Message Passing Interface
- Defines a standard for programs to communicate with each other
- Describes a set of functions and their expected behavior
- MPI\_Send() is described but not *implemented*
- Implementation is left to others

## 3.2 Blocking Send and Receive Operations

### 3.2.1 Blocking Send

The syntax of the **blocking send** procedure is given below.

MPI\_SEND(buf, count, datatype, dest, tag, comm)

IN	buf	initial address of send buffer (choice)
IN	count	number of elements in send buffer (non-negative integer)
IN	datatype	datatype of each send buffer element (handle)
IN	dest	rank of destination (integer)
IN	tag	message tag (integer)
IN	comm	communicator (handle)

#### C binding

```
int MPI_Send(const void *buf, int count, MPI_Datatype datatype, int dest,  
            int tag, MPI_Comm comm)
```

```
int MPI_Send_c(const void *buf, MPI_Count count, MPI_Datatype datatype,  
              int dest, int tag, MPI_Comm comm)
```

Official description for MPI\_Send

# What is MPI?

- There are many MPI implementations
- Each implementation must be compliant with MPI description
- On Discovery/Endeavour you will find
  - openmpi
  - mpich
  - intel-mpi
  - mvapich
- Each solves the same problems but in different ways

```
int MPI_Send(const void *buf, int count, MPI_Datatype type, int dest,
             int tag, MPI_Comm comm)
{
    int rc = MPI_SUCCESS;

    SPC_RECORD(OMPI_SPC_SEND, 1);

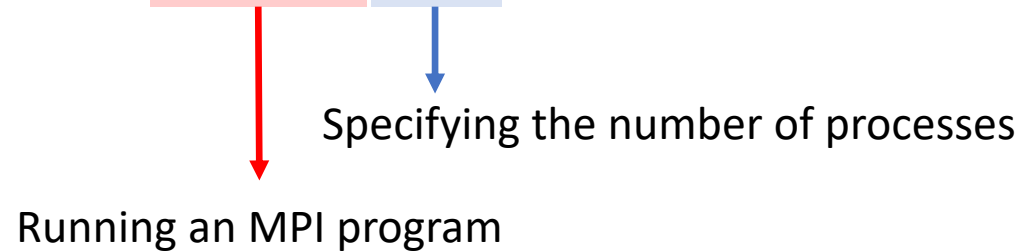
    MEMCHECKER(
        memchecker_datatype(type);
        memchecker_call(&opal_memchecker_base_isdefined, buf, count, type);
        memchecker_comm(comm);
    );

    if ( MPI_PARAM_CHECK ) {
        OMPI_ERR_INIT_FINALIZE(FUNC_NAME);
        if (ompi_comm_invalid(comm)) {
            return OMPI_ERRHANDLER_NOHANDLE_INVOKE(MPI_ERR_COMM, FUNC_NAME);
        } else if (count < 0) {
            rc = MPI_ERR_COUNT;
        } else if (tag < 0 || tag > mca_pml.pml_max_tag) {
            rc = MPI_ERR_TAG;
        } else if (ompi_comm_peer_invalid(comm, dest) &&
                    (MPI_PROC_NULL != dest)) {
            rc = MPI_ERR_RANK;
        } else {
            OMPI_CHECK_DATATYPE_FOR_SEND(rc, type, count);
            OMPI_CHECK_USER_BUFFER(rc, buf, type, count);
        }
        OMPI_ERRHANDLER_CHECK(rc, comm, rc, FUNC_NAME);
    }
}
```

OpenMPI version of MPI\_Send

Serial `echo Hello World!`

Parallel `mpirun -n 4 echo Hello World!`



Running an MPI program

Specifying the number of processes

The name “`mpirun`” is not part of the standard, other names include:

SLURM: `srun`

Mpich2: `mpiexec`

IBM SP: `poe`

Stampede2: `ibrun`

If you don’t have enough CPUs to run this command on your machine, you may get:

`There are not enough slots available in the system`

Solution: `mpirun --oversubscribe -n 4 echo Hello World!`



Serial `echo` Hello World!

Parallel `mpirun` `-n 4` `echo` Hello World!



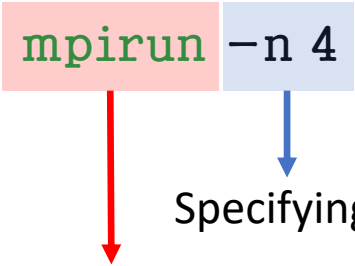
Specifying the number of processes

Running an MPI program

- MPI copies a program several times and run them individually
- How to tell a **program**, it is being executed by an MPI command and all pieces should work together?
  - **In C and Fortran:**  
Start the program with function: `MPI_INIT` and clean up with: `MPI_FINALIZE`
  - **In Python:**  
Handled by the library: `from mpi4py import MPI`

Serial `echo` Hello World!

Parallel `mpirun -n 4 echo` Hello World!

  
Running an MPI program

- MPI copies a program several times and run them individually
- How to tell a **program**, it is being executed by an MPI command and all pieces should work together?
  - **In C and Fortran:**  
Start the program with function: `MPI_INIT` and clean up with: `MPI_FINALIZE`
  - **In Python:**  
Handled by the library: `from mpi4py import MPI`

- `np = MPI.COMM_WORLD.Get_size()` Total number of processes
- `myrank = MPI.COMM_WORLD.Get_rank()` Rank of each process (between 0,np-1)

# Hello World MPI

```
from mpi4py import MPI

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

print(f"Hello from rank {rank}")
```

- Each process will:
  - Start up
  - Find out their "rank" number
  - Print rank to screen
- See [examples/hello\\_world](#) for code

# Hello World MPI

```
from mpi4py import MPI

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

print(f"Hello from rank {rank}")
```

← Communicator - Group of processes

← ID number within the communicator

- Each process will:
  - Start up
  - Find out their "rank" number
  - Print rank to screen
- See [examples/hello\\_world](#) for code

# Getting Compute resources

- MPI programs must be run in Slurm jobs
- Each "rank" must be given a Slurm "task"
- Use reservation for this workshop:
- `#SBATCH --reservation=bootcamp`
- `#SBATCH --account=hpcsuppt_613`
- `#SBATCH --partition=gpu`
- We will be using mpi4py and numpy
- If not installed, use `pip install mpi4py numpy`

# Run from interactive job

```
$ salloc --ntasks=4 --reservation=bootcamp  
--account=hpcsuppt_613 --time=1:00:00 --partition=gpu
```

```
$ module load usc
```

```
$ module load python
```

```
$ srun python3 hello_world_mpi.py
```

```
Hello from rank 0!
```

```
Hello from rank 2!
```

```
Hello from rank 3!
```

```
Hello from rank 1!
```

# Run from batch job

```
#!/bin/bash
#SBATCH --ntasks=4
#SBATCH --partition=debug
#SBATCH --mem-per-cpu=2GB

module load usc
module load python

srun python3 hello_world_mpi.py
```

- Use interactive job to help build job script
- Everything we did in interactive job
- Output saved to slurm-xxxxx.out file

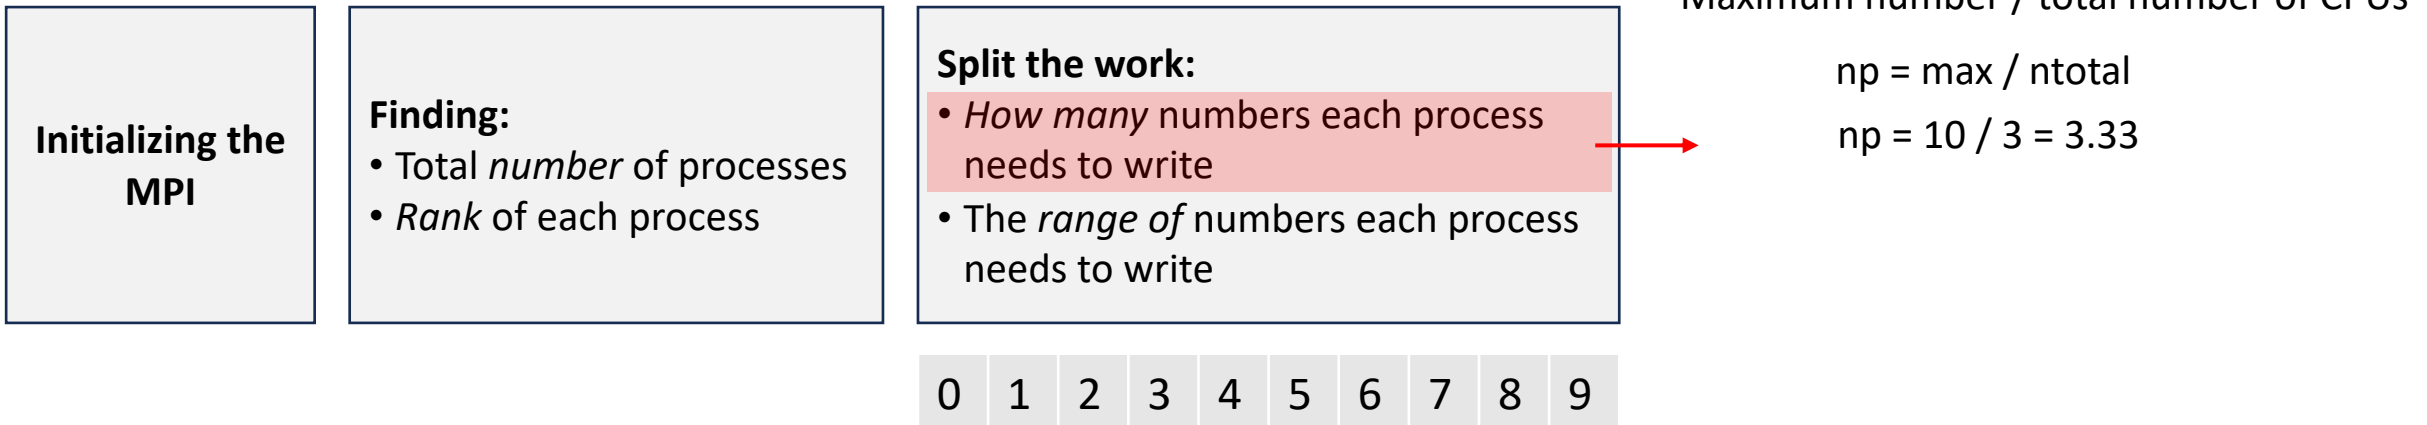
# Doing Calculations in Parallel

**Serial**

Print from 0 to 9

**Parallel**

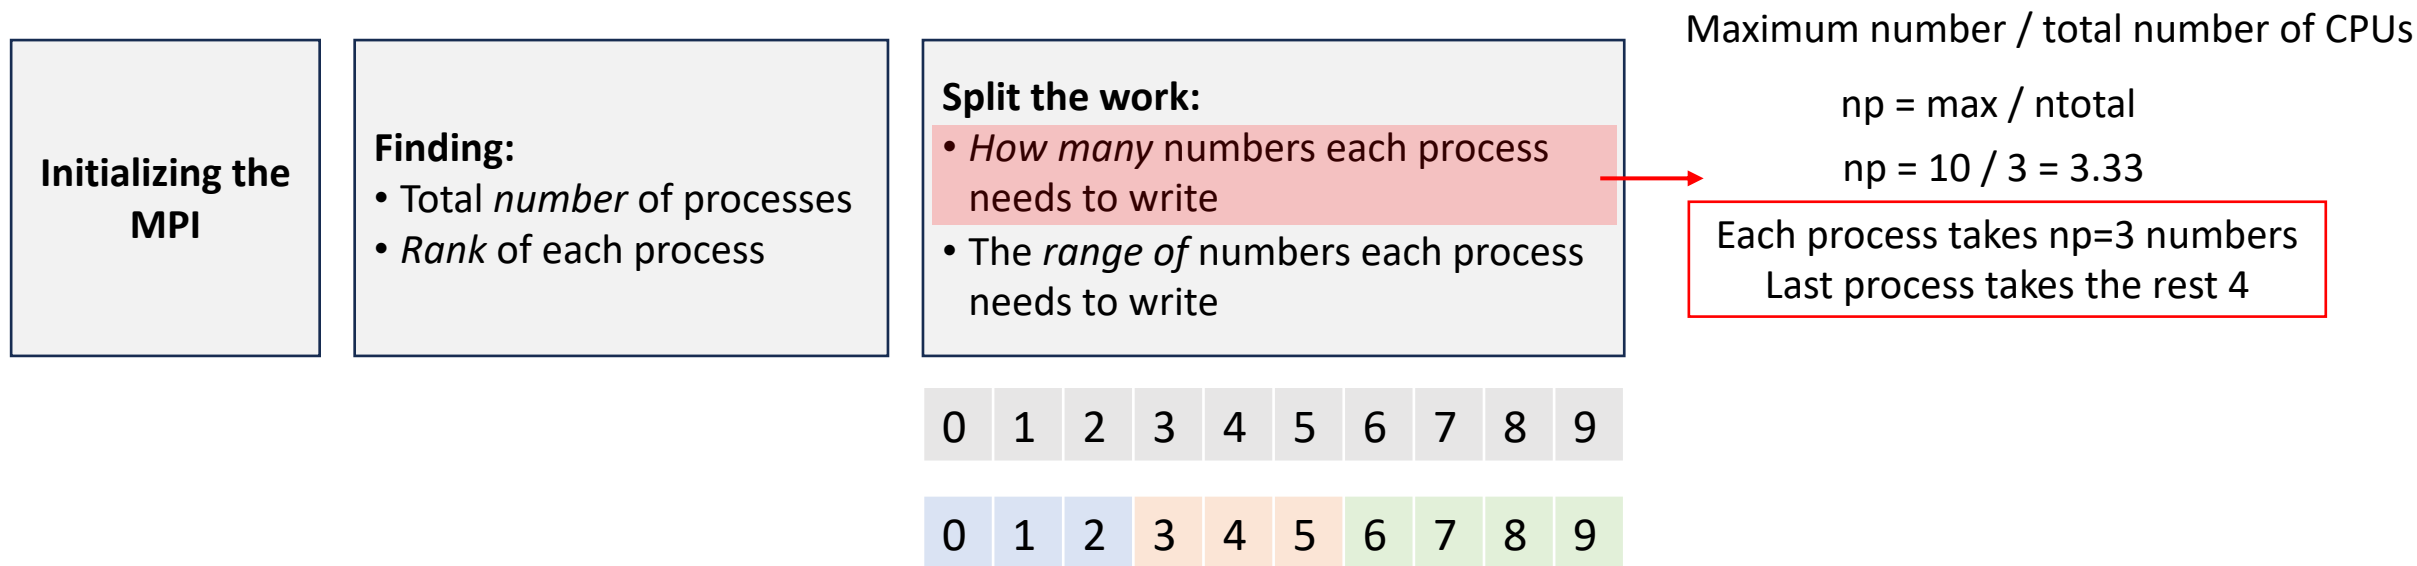
Print from 0 to 9 with 3 processes, each process prints “its rank”, and “its portion of the numbers”  
Last process carries whatever is left





# Doing Calculations in Parallel

- Serial
  - Print from 0 to 9
- Parallel
  - Print from 0 to 9 with 3 processes, each process prints “its rank”, and “its portion of the numbers”
  - Last process carries whatever is left



# Doing Calculations in Parallel

Serial

Print from 0 to 9

Parallel

Print from 0 to 9 with 3 processes, each process prints “its rank”, and “its portion of the numbers”  
Last process carries whatever is left

Initializing the  
MPI


**Finding:**

- Total *number* of processes
- *Rank* of each process

**Split the work:**

- *How many* numbers each process needs to write
- The *range of* numbers each process needs to write

Each process takes **np=3** numbers  
Last process takes the rest 4

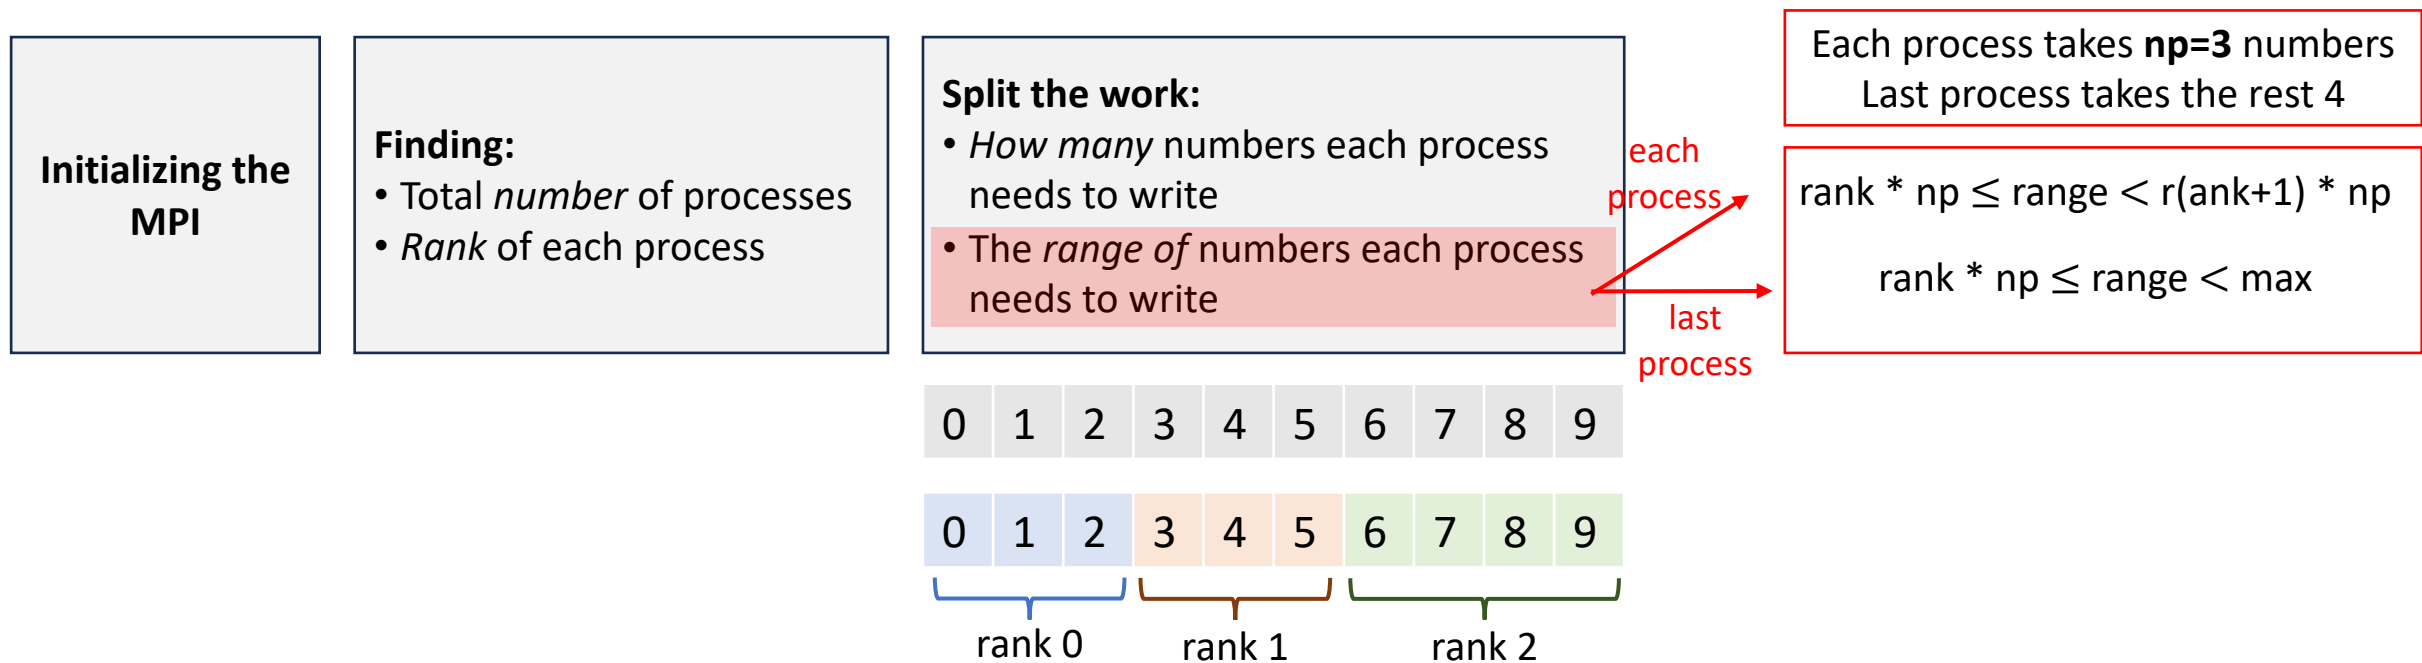


# Doing Calculations in Parallel

- Serial

Print from 0 to 9
- Parallel

Print from 0 to 9 with 3 processes, each process prints “its rank”, and “its portion of the numbers”  
Last process carries whatever is left



# Doing Calculations in Parallel

Serial	Print from 0 to 9
Parallel	Print from 0 to 9 with 3 processes, each process prints “its rank”, and “its portion of the numbers” Last process carries whatever is left

```
from mpi4py import MPI
```

Initializing the MPI

# Doing Calculations in Parallel

Serial	Print from 0 to 9
Parallel	Print from 0 to 9 with 3 processes, each process prints “its rank”, and “its portion of the numbers” Last process carries whatever is left

```
from mpi4py import MPI
```

```
max=10
```

```
ntotal = MPI.COMM_WORLD.Get_size()
```

```
myrank = MPI.COMM_WORLD.Get_rank()
```

## Initializing the MPI

### Finding:

- Total *number* of processes
- *Rank* of each process

# Doing Calculations in Parallel

Serial	Print from 0 to 9
Parallel	Print from 0 to 9 with 3 processes, each process prints “its rank”, and “its portion of the numbers” Last process carries whatever is left

```
from mpi4py import MPI
```

```
max=10
```

```
ntotal = MPI.COMM_WORLD.Get_size()
```

```
myrank = MPI.COMM_WORLD.Get_rank()
```

```
np = max//ntotal #gives the quotient  
remainder=max % ntotal
```

## Initializing the MPI

### Finding:

- Total *number* of processes
- *Rank* of each process

### Split the work:

- *How many* numbers each process needs to write

# Doing Calculations in Parallel

Serial	Print from 0 to 9
Parallel	Print from 0 to 9 with 3 processes, each process prints “its rank”, and “its portion of the numbers” Last process carries whatever is left

```
from mpi4py import MPI

max=10

ntotal = MPI.COMM_WORLD.Get_size()
myrank = MPI.COMM_WORLD.Get_rank()

np = max//ntotal #gives the quotient
remainder=max % ntotal

if (myrank is ntotal-1):
    print ('myrank=', myrank, ' start=', myrank*np, ' end=', (myrank+1)*np+remainder)
else:
    print ('myrank=', myrank, ' start=', myrank*np, ' end=', (myrank+1)*np)
```

## Initializing the MPI

### Finding:

- Total *number* of processes
- *Rank* of each process


### Split the work:

- *How many* numbers each process needs to write

- The *range* of numbers each process needs to write

# Point to Point Communication - Send


- Simplest way to send data
- `comm.send(buf, dest, tag)`
- Different parts
  - `buf` – the thing you want to send
  - `dest` – rank to send to
  - `tag` – label for data
- Send will wait for successful receive before continuing





# Point to Point Communication - Recv

- Simplest way to receive data
- `comm.recv(buf, source=ANY_SOURCE, tag=ANY_TAG)`
- Different parts
  - `buf` – container for data
  - `source` – rank data is expected from
  - `tag` – label for data
- Recv will wait for data before continuing
- Recv will only accept data from `dest` with matching `tag`



# Send/Recv example

```
from mpi4py import MPI

comm = MPI.COMM_WORLD
rank=comm.Get_rank()

print(f'Rank {rank} starting up ...')

good_tag=7
data = None

if rank == 0:
    print(f'Rank {rank}: Sending data to rank 1 with tag {good_tag}')
    comm.send(data,dest=1,tag=good_tag)

if rank == 1:
    #data=None
    print(f'Rank {rank}: Waiting for data from rank 0 with tag {good_tag}')
    data=comm.recv(source=0, tag=good_tag)
    print(f'Rank {rank}: Got data: {data} from rank 0 with tag {good_tag}')

print(f'Rank {rank} shutting down with data={data}.')
```

# Send/Recv example - Initialize

```
from mpi4py import MPI

comm = MPI.COMM_WORLD
rank=comm.Get_rank()

print(f'Rank {rank} starting up ...')

good_tag=7
data = None

if rank == 0:
    data=412
    print(f'Rank {rank}: Sending data to rank 1 with tag {good_tag}')
    comm.send(data,dest=1,tag=good_tag)

if rank == 1:
    #data=None
    print(f'Rank {rank}: Waiting for data from rank 0 with tag {good_tag}')
    data=comm.recv(source=0, tag=good_tag)
    print(f'Rank {rank}: Got data: {data} from rank 0 with tag {good_tag}')

print(f'Rank {rank} shutting down with data={data}.')
```

# Send/Recv example - Rank 0

```
from mpi4py import MPI

comm = MPI.COMM_WORLD
rank=comm.Get_rank()

print(f'Rank {rank} starting up ...')

good_tag=7
data = None

if rank == 0:
    data=412
    print(f'Rank {rank}: Sending data to rank 1 with tag {good_tag}')
    comm.send(data,dest=1,tag=good_tag)

if rank == 1:
    #data=None
    print(f'Rank {rank}: Waiting for data from rank 1 with tag {good_tag}')
    data=comm.recv(source=0, tag=good_tag)
    print(f'Rank {rank}: Got data: {data} from rank 0 with tag {good_tag}')

print(f'Rank {rank} shutting down with data={data}.')
```

# Send/Recv example - Rank 1

```
from mpi4py import MPI

comm = MPI.COMM_WORLD
rank=comm.Get_rank()

print(f'Rank {rank} starting up ...')

good_tag=7
data = None

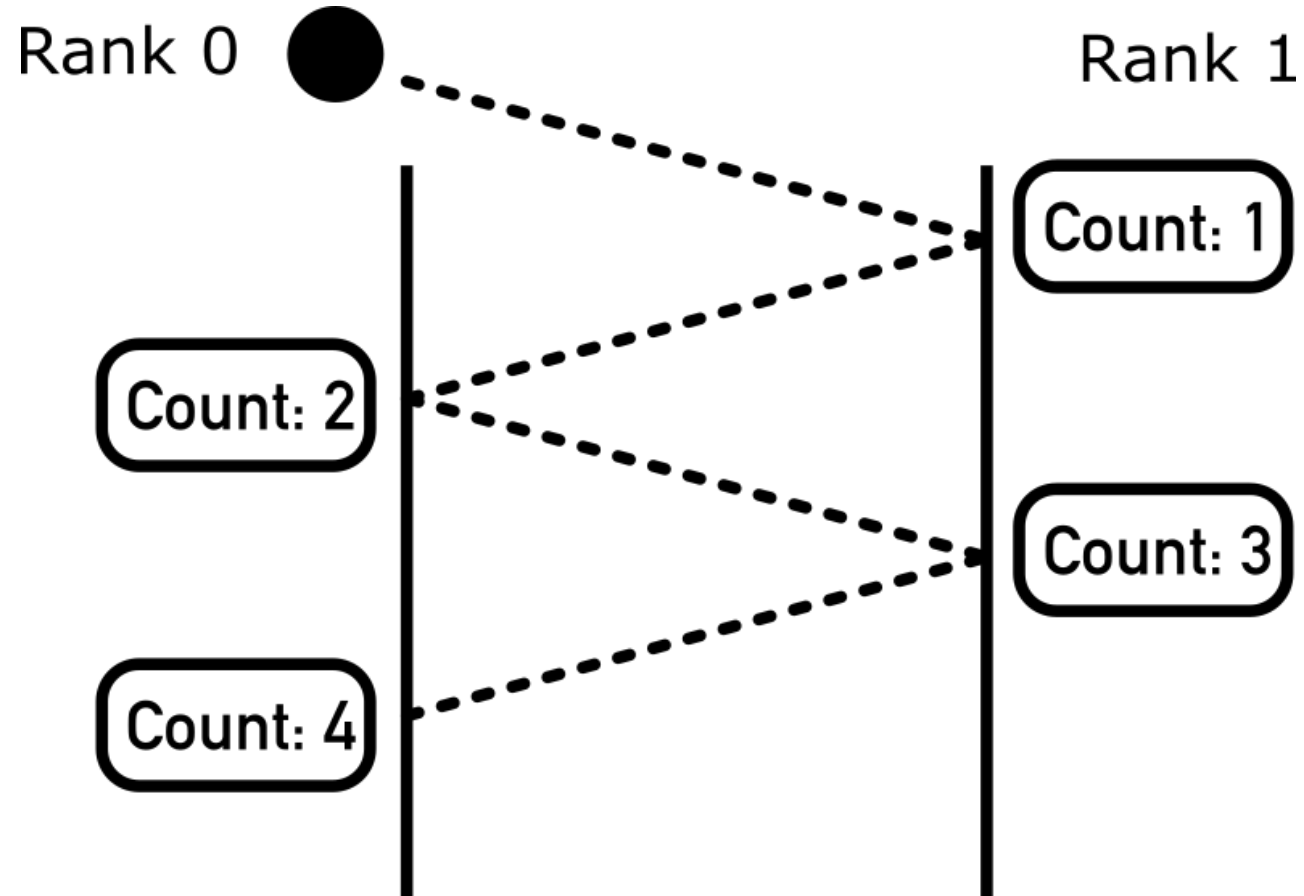
if rank == 0:
    data=412
    print(f'Rank {rank}: Sending data to rank 1 with tag {good_tag}')
    comm.send(data,dest=1,tag=good_tag)

if rank == 1:
    #data=None
    print(f'Rank {rank}: Waiting for data from rank 1 with tag {good_tag}')
    data=comm.recv(source=0, tag=good_tag)
    print(f'Rank {rank}: Got data: {data} from rank 0 with tag {good_tag}')

print(f'Rank {rank} shutting down with data={data}.')
```

# MPI Ping Pong

- Classic MPI example
- Two ranks alternate sending a message
- Counter is incremented on receipt
- Stop after a N rounds



# MPI Ping Pong - Code

```
from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD
rank = comm.Get_rank()
world_size = MPI.COMM_WORLD.Get_size()

if world_size != 2:
    print("Only two can play")
    sys.exit(1)

print(f'Rank {rank} starting up...')

counter = 0
max_counter = 10

if rank == 0:
    partner = 1
if rank == 1:
    partner = 0

while counter < max_counter:
    counter=comm.recv(source=partner)
    print(f'Rank {rank}: Got message {counter} from rank {partner}')
    counter = counter + 1

    print(f'Rank {rank}: sending message {counter} to rank {partner}')
    comm.send(counter,dest=partner)
```

# MPI Ping Pong - Initialize

```
from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD
rank = comm.Get_rank()
world_size = MPI.COMM_WORLD.Get_size()

if world_size != 2:
    print("Only two can play")
    sys.exit(1)

print(f'Rank {rank} starting up...')

if rank == 0:
    partner = 1
if rank == 1:
    partner = 0
```



# MPI Ping Pong - Send/Receive loop

```
counter = 0
max_counter = 10

while counter < max_counter:
    counter=comm.recv(source=partner)
    print(f'Rank {rank}: Got message {counter} from rank {partner}')
    counter = counter + 1


    print(f'Rank {rank}: sending message {counter} to rank {partner}')
    comm.send(counter,dest=partner)
```

# MPI Ping Pong

- Something is wrong with this example!
- Run the program on your own
- What could you change to make it work?

# Deadlock

- Both ranks are "receiving"
- Can't proceed until done




```
counter = 0
max_counter = 10
```

```
while counter < max_counter:
    counter=comm.recv(source=partner)
    print(f'Rank {rank}: Got message {counter} from rank {partner}')
    counter = counter + 1

    print(f'Rank {rank}: sending message {counter} to rank {partner}')
    comm.send(counter,dest=partner)
```

# Deadlock

- Be careful with message coordination



```
counter = 0
max_counter = 10

if rank == 0:
    comm.send(counter, dest=partner)

while counter < max_counter:
    counter=comm.recv(source=partner)
    print(f'Rank {rank}: Got message {counter} from rank {partner}')

    print(f'Rank {rank}: sending message {counter} to rank {partner}')
    comm.send(counter, dest=partner)
    counter = counter + 1
```

# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

0			0		0		
						0	0
	0						
	0						
		0				0	
			0				

# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

Initializing  
the MPI

**Finding:**

- Total *number* of processes
- *Rank* of each process

0			0		0		
						0	0
	0						
	0						
		0				0	
			0				

# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

Initializing  
the MPI

Finding:

- Total *number* of processes
- *Rank* of each process

Creating  
the matrix

Split the work:

- The *range of* elements each process needs to read
- Calculate

0			0		0		
						0	0
	0						
	0						
		0				0	
			0				

# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

<b>Initializing the MPI</b>	<b>Finding:</b> <ul style="list-style-type: none"><li>• Total <i>number</i> of processes</li><li>• <i>Rank</i> of each process</li></ul>	<b>Creating the matrix</b>	<b>Split the work:</b> <ul style="list-style-type: none"><li>• The <i>range of</i> elements each process needs to read</li><li>• Calculate</li></ul>
-----------------------------	--	----------------------------	--

0			0		0		
						0	0
	0						
	0						
		0				0	
			0				



# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

Initializing  
the MPI

Finding:

- Total *number* of processes
- *Rank* of each process

Creating  
the matrix

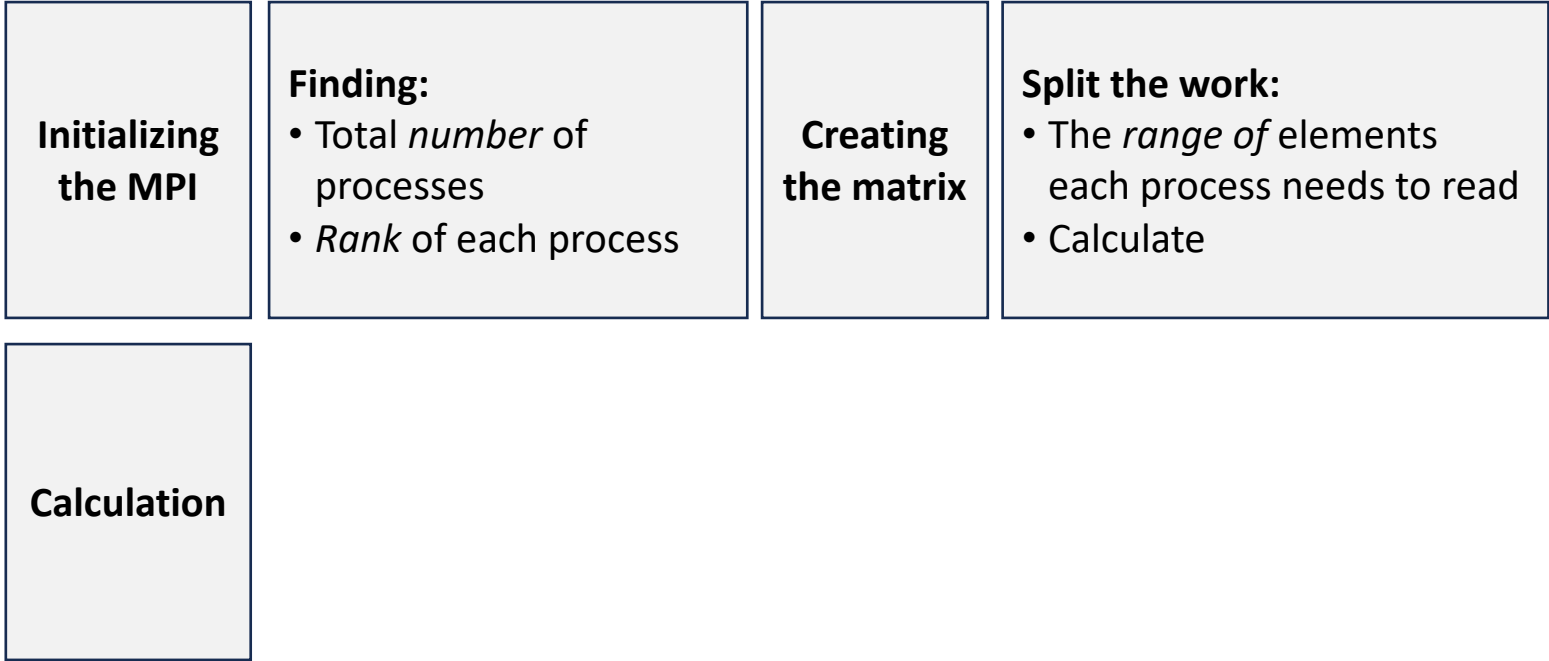
Split the work:

- The *range of* elements each process needs to read
- Calculate

0			0		0		
rank0:				rank1:			
						0	0
	0						
	0						
rank2:				rank3:			
			0				

# Example


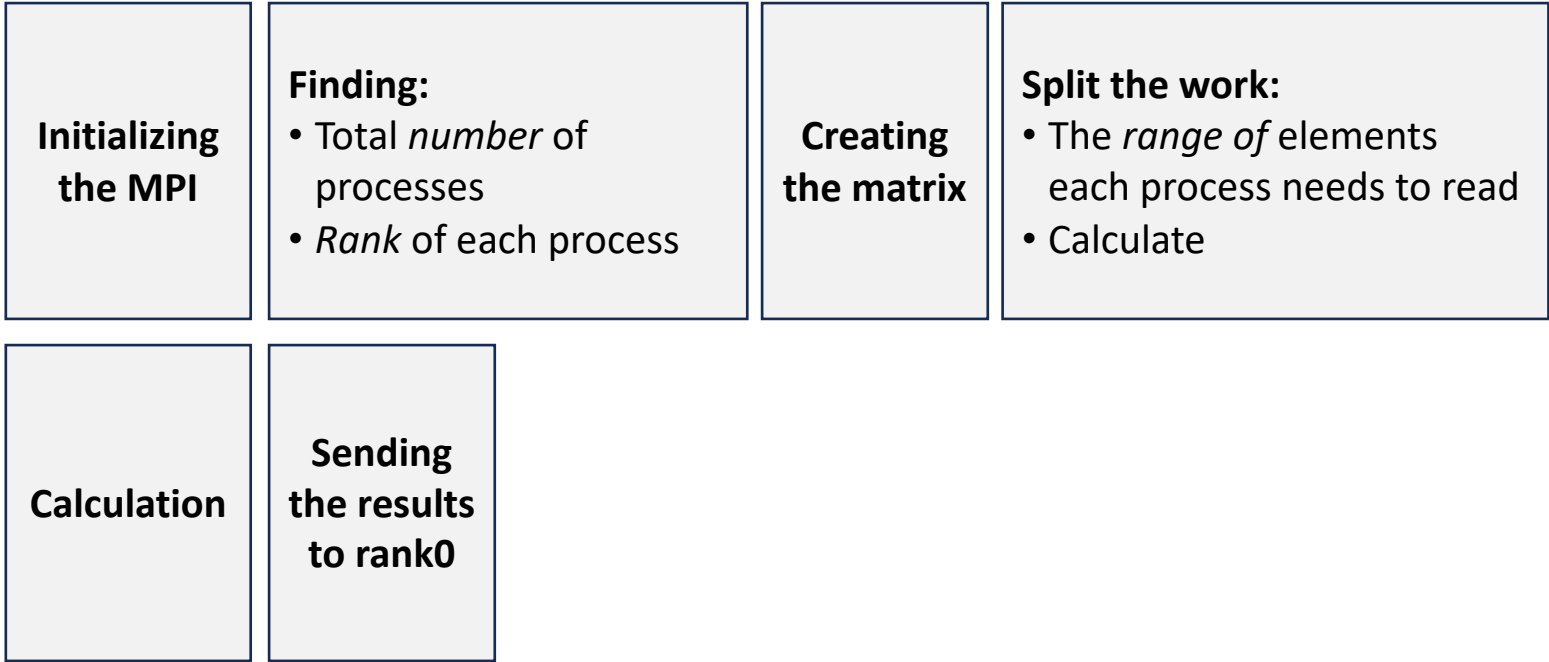
Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results



0			0		0		
	rank0:			rank1:			
	nz=3			nz=3	0	0	
	0						
	0						
	rank2:			rank3:	0		
	nz=3			nz=1			
			0				


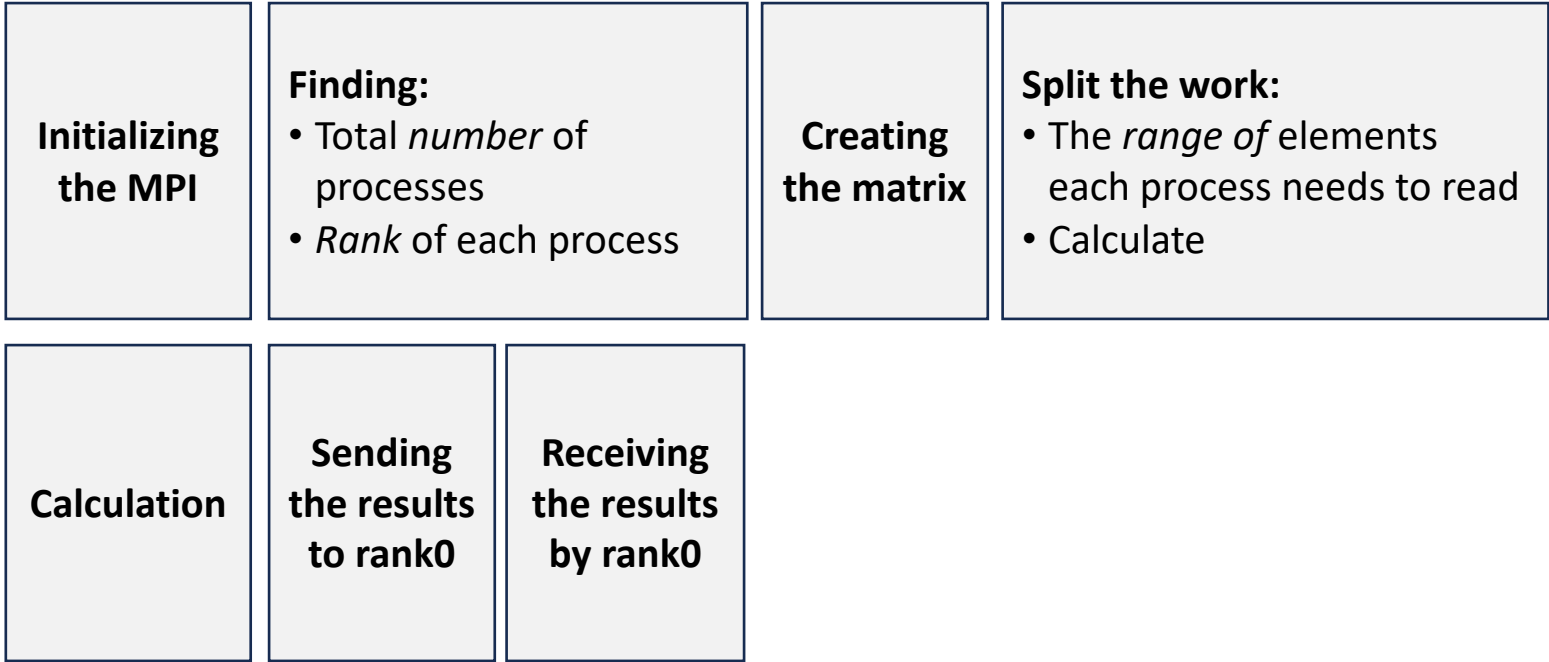
# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results



# Example


Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results



# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

<b>Initializing the MPI</b>	<b>Finding:</b> <ul style="list-style-type: none"><li>• Total <i>number</i> of processes</li><li>• <i>Rank</i> of each process</li></ul>	<b>Creating the matrix</b>	<b>Split the work:</b> <ul style="list-style-type: none"><li>• The <i>range of</i> elements each process needs to read</li><li>• Calculate</li></ul>	
<b>Calculation</b>	<b>Sending the results to rank0</b>	<b>Receiving the results by rank0</b>	<b>Calculation in rank0</b>	<b>Printing the results by rank0</b>



## Example

## Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

```
from mpi4py import MPI
```


## Initializing the MPI

```
comm = MPI.COMM_WORLD
ntotal = MPI.COMM_WORLD.Get_size()
rank = MPI.COMM_WORLD.Get_rank()
```

### Finding:


- Total *number* of processes
- *Rank* of each process

```
import numpy as np
np.random.seed(10)
A=np.random.rand(8,8)
i1ist=np.random.randint(0,7,10)
j1ist=np.random.randint(0,7,10)
A[i1ist,j1ist]=0
```



## Creating the matrix

```
npi=int(2)
npj=int(ntotal/npi)
```



### Split the work:

- *How many* elements each process needs to write

# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

```
from mpi4py import MPI
```


Initializing the MPI

```
comm = MPI.COMM_WORLD
ntotal = MPI.COMM_WORLD.Get_size()
rank = MPI.COMM_WORLD.Get_rank()
```

Finding:

- Total *number* of processes
- *Rank* of each process


```
import numpy as np
np.random.seed(10)
A=np.random.rand(8,8)
ilist=np.random.randint(0,7,10)
jlist=np.random.randint(0,7,10)
A[ilist,jlist]=0
```



Creating the matrix

```
npi=int(2)
npj=int(ntotal/npi)

rank_i=rank//npi
rank_j=rank%npi
```



Split the work:

- *How many* elements each process needs to write

# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

```
from mpi4py import MPI
```


Initializing the MPI

```
comm = MPI.COMM_WORLD
ntotal = MPI.COMM_WORLD.Get_size()
rank = MPI.COMM_WORLD.Get_rank()
```

Finding:

- Total *number* of processes
- *Rank* of each process

```
import numpy as np
np.random.seed(10)
A=np.random.rand(8,8)
ilist=np.random.randint(0,7,10)
jlist=np.random.randint(0,7,10)
A[ilist,jlist]=0
```




Creating the matrix

```
npi=int(2)
npj=int(ntotal/npi)

rank_i=rank//npi
rank_j=rank%npi

ni,nj=A.shape
chunk_size_i=int(ni/npi)
chunk_size_j=int(nj/npj)
```



Split the work:

- *How many* elements each process needs to write




# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

```
rank_i=rank//npi
rank_j=rank%npi
```

```
ni,nj=A.shape
chunk_size_i=int(ni/npi)
chunk_size_j=int(nj/npj)
```

```
A_tmp=A[rank_i*chunk_size_i:(rank_i+1)*chunk_size_i,
        rank_j*chunk_size_j:(rank_j+1)*chunk_size_j]
```



## Split the work:

- *How many* elements each process needs to write


# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

```
rank_i=rank//npi
rank_j=rank%npi
```

```
ni,nj=A.shape
chunk_size_i=int(ni/npi)
chunk_size_j=int(nj/npj)
```

```
A_tmp=A[rank_i*chunk_size_i:(rank_i+1)*chunk_size_i,
        rank_j*chunk_size_j:(rank_j+1)*chunk_size_j]
```



**Split the work:**

- *How many* elements each process needs to write

```
nz_local =len(np.where(A_tmp==0)[0])
```

**Calculation**

# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

```
nz_local =len(np.where(A_tmp==0)[0])
```


**Calculation**

# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

```
nz_local = len(np.where(A_tmp==0)[0])
```

```
if rank != 0:  
    comm.send(nz_local, dest=0, tag=7)
```



Calculation


Sending the results to  
rank0

# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

```
nz_local =len(np.where(A_tmp==0)[0])
```


```
if rank !=0:  
    comm.send(nz_local,dest=0,tag=7)
```



**Calculation**

**Sending the results to  
rank0**

```
if rank==0:  
    nz_list=np.zeros((ntotal,),int)  
  
    nz_list[0]=nz_local  
  
    for i in range(1,ntotal):  
        nz_list[i]=comm.recv(source=i,tag=7)
```



**Receiving the results by  
rank0**


# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

```
nz_local = len(np.where(A_tmp==0)[0])
```


Calculation

```
if rank !=0:  
    comm.send(nz_local,dest=0,tag=7)
```



Sending the results to rank0

```
if rank==0:  
    nz_list=np.zeros((ntotal,),int)  
  
    nz_list[0]=nz_local  
  
    for i in range(1,ntotal):  
        nz_list[i]=comm.recv(source=i,tag=7)
```



Receiving the results by rank0

```
if rank==0:  
    nz_global=sum(nz_list)
```

Calculation


```
if rank==0:  
    print(nz_global)
```

Printing the results

# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results  
What if we had more processes?


## 4 Processes




# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results  
What if we had more processes?

**4 Processes**



**8 Processes**







# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results  
What if we had more processes?


**4 Processes**



**8 Processes**




**16 Processes**




# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results  
What if we had more processes?


4 Processes



8 Processes




16 Processes



16 Processes:  
Smarter implementation


**Collective Communication**




# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results  
What if we had more processes?


4 Processes



8 Processes




16 Processes



16 Processes:

Smarter implementation

**Collective Communication**



This specific case is called:

**MPI\_Gather**


# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

```
nz_local =len(np.where(A_tmp==0)[0])
```


Calculation

```
if rank !=0:  
    comm.send(nz_local,dest=0,tag=7)
```



Sending the results to rank0

```
if rank==0:  
    nz_list=np.zeros((ntotal,),int)  
  
    nz_list[0]=nz_local  
  
    for i in range(1,ntotal):  
        nz_list[i]=comm.recv(source=i,tag=7)
```



Receiving the results by rank0

```
if rank==0:  
    nz_global=sum(nz_list)
```

Calculation

```
if rank==0:  
    print(nz_global)
```

Printing the results


# Example

Count the number of zero elements of a matrix in parallel (with 4 processes) and print the results

```
nz_local = len(np.where(A_tmp==0)[0])
```

Calculation

```
nz_list = comm.gather(nz_local, root=0)
```



Gather `nz_local`

```
if rank==0:  
    nz_global=sum(nz_list)
```


Calculation

```
if rank==0:  
    print(nz_global)
```

Printing the results

# Collective Communication

## Gather




At rank\_i: value=a\_i

```
vector = comm.gather(value, root=0)
```

At rank\_0: vector = [a0, a1, a2, ..., a15]


# Collective Communication

Gather



At rank_i:	value=a_i
vector = comm.gather( <b>value</b> , root=0)	
At rank_0:	vector = [a0, a1, a2, ..., a15]


Scatter



At rank_0:	vector = [a0, a1, a2, ..., a15]
value = comm.scatter( <b>vector</b> , root=0)	
At rank_0:	value=a_i


# Collective Communication

**Gather**




At rank_i:	value=a_i
vector = comm.gather( <b>value</b> , root=0)	
At rank_0:	vector = [a0, a1, a2, ..., a15]

**Scatter**



At rank_0:	vector = [a0, a1, a2, ..., a15]
value = comm.scatter( <b>vector</b> , root=0)	
At rank_i:	value=a_i

**Broadcast**




At rank_0:	value=a_0
value = comm.broadcast( <b>value</b> , root=0)	
At rank_i:	value=a_0




# Collective Communication

## Gather



At rank_i:	value=a_i
vector = comm.gather(value, root=0)	
At rank_0:	vector = [a0, a1, a2, ..., a15]

## Allgather



At rank_i:	value=a_i
vector = comm.allgather(value)	
At rank_i:	vector=[a0, a1, a2, ..., a15]

## Type of collective operations

### Data Movement

- Broadcast
- Scatter/Gather
- Allgather/AlltoAll

### Computation

- Reduce
- Allreduce
  - SUM
  - MIN/MAX
  - ...
- Scan

### Synchronization

- Barrier

# Thank you!

- Questions?
- [carc-support@usc.edu](mailto:carc-support@usc.edu)
- Office hours: LVL 3L and Zoom 2:30-5pm Every Tuesday
- Thanks to Marco for help on these slides!